

Lab1-Backpropagation Report

313551104 黃暉洺

1. Introduction (20%)

本次實作了兩層 hidden layer 的神經網路，只使用 numpy 完成全部的運算，程式包含 layer 及 MLP 兩個 class，和其他數學運算的副程式，layer 可執行 feedforward 和 backpropagation，運用 linear 跟 XOR 兩個數據集做訓練與測試，最後比較各種參數之間在效能上的差異。

2. Experiment setups (30%) :

A. Sigmoid functions

```
def sigmoid(x):  
    return 1.0 / (1.0 + np.exp(-x))
```

B. Neural network

```
class MLP:  
    def __init__(self, hidden_size = 5, learning_rate = 0.5, activate = "none"):  
        self.learning_rate = learning_rate  
        self.hidden1 = layer(2, hidden_size, activate)  
        self.hidden2 = layer(hidden_size, hidden_size, activate)  
        self.output = layer(hidden_size, 1, activate)  
        self.loss = []  
  
    def train(self, x, ground_truth, epoch = 100000):  
        plt.figure(figsize = (30, 30))  
        plt.subplot(2, 1, 1)  
        plt.title('Learning Curves', fontsize = 30)  
        plt.xlabel('epoch', fontsize = 30)  
        plt.ylabel('loss', fontsize = 30)  
        for i in range(epoch):  
            self.hidden1.forward(x)  
            self.hidden2.forward(self.hidden1.z)  
            self.output.forward(self.hidden2.z)  
  
            loss, gradient = self.cost(ground_truth)  
            self.loss.append(loss)  
            if not i % 5000:  
                print("epoch ", i, " loss : ", loss)  
  
            self.output.back(gradient, self.learning_rate)  
            self.hidden2.back(self.output.gradient, self.learning_rate)  
            self.hidden1.back(self.hidden2.gradient, self.learning_rate)  
        plt.plot(self.loss)  
  
        show_result(x, ground_truth, self.output.z)  
        for i in range(ground_truth.size):  
            print("Iter", i + 1, " | Ground truth: ", ground_truth[i], " | prediction: ", self.output.z[i], "  
        print("loss=", loss, " accuracy=", 100 * sum((self.output.z > 0.5) == (ground_truth == 1)) / ground_truth  
  
    def cost(self, y_hat):  
        return MSE(self.output.z, y_hat), y_grad(self.output.z, y_hat)  
  
    def test(self, x, ground_truth):  
        self.hidden1.forward(x)  
        self.hidden2.forward(self.hidden1.z)  
        self.output.forward(self.hidden2.z)  
        loss, grad = self.cost(ground_truth)  
        print(" ===== testing set ===== ")  
        show_result(x, ground_truth, self.output.z)  
        for i in range(ground_truth.size):  
            print("Iter", i + 1, " | Ground truth: ", ground_truth[i], " | prediction: ", self.output.z[i], "  
        print("loss=", loss, " accuracy=", 100 * sum((self.output.z > 0.5) == (ground_truth == 1)) / ground_truth
```

C. Backpropagation

```

class layer:
    def __init__(self, input_size, output_size, activate = "none"):
        self.input_size = input_size
        self.output_size = output_size
        self.activate = activate
        self.gradient = 0
        self.w = np.random.normal(-1, 1, (input_size, output_size))
        self.b = np.random.normal(-1, 1, (1, output_size))
        self.a = []
        self.z = []

    def forward(self, x):
        self.x = x
        self.a = np.dot(x, self.w) + self.b
        if(self.activate == "sigmoid"):
            self.z = sigmoid(self.a)
        elif(self.activate == "ReLU"):
            self.z = ReLU(self.a)
        else:
            self.z = self.a


    def back(self, gradient, learning_rate):
        if self.activate == "sigmoid":
            gradient *= derivative_sigmoid(self.z)
        elif self.activate == "ReLU":
            gradient *= derivative_ReLU(self.z)
        else:
            pass


        self.w -= learning_rate * np.dot(self.x.T, gradient)
        self.b -= learning_rate * np.sum(gradient, axis = 0)
        self.gradient = np.dot(gradient, self.w.T)

```

3. Results of your testing (20%)

A. Screenshot and comparison figure





B. Show the accuracy of your prediction

Linear :

```
x1, y1 = generate_linear()
model = MLP(hidden_size = 5, activate = "sigmoid", learning_rate = 0.1)
model.train(x1, y1, epoch = 100000)
```

✓ 4.4s

```
epoch 0   loss : 0.3222392885432805
epoch 5000 loss : 0.02616356246458329
epoch 10000 loss : 0.01174779135501703
epoch 15000 loss : 0.008510572394455846
epoch 20000 loss : 0.00692084150529909
epoch 25000 loss : 0.005909300050339386
epoch 30000 loss : 0.00517039997620771
epoch 35000 loss : 0.004586251604458222
epoch 40000 loss : 0.004103229797072347
epoch 45000 loss : 0.003693583068616542
epoch 50000 loss : 0.003341050109749299
epoch 55000 loss : 0.0030349077714084778
epoch 60000 loss : 0.002767363823042213
epoch 65000 loss : 0.0025323640557419116
epoch 70000 loss : 0.002325015394309402
epoch 75000 loss : 0.002141281055711663
epoch 80000 loss : 0.0019777979312595203
epoch 85000 loss : 0.001831751416627836
epoch 90000 loss : 0.0017007801040933124
epoch 95000 loss : 0.0015828986711753186
```

```

Iter 1 | Ground truth: [1] | prediction: [0.9881353] |
Iter 2 | Ground truth: [0] | prediction: [4.49336206e-06] |
Iter 3 | Ground truth: [1] | prediction: [0.9999837] |
Iter 4 | Ground truth: [1] | prediction: [0.99998683] |
Iter 5 | Ground truth: [0] | prediction: [4.95955647e-06] |
Iter 6 | Ground truth: [1] | prediction: [0.99998707] |
Iter 7 | Ground truth: [1] | prediction: [0.99998807] |
Iter 8 | Ground truth: [0] | prediction: [0.00228932] |
Iter 9 | Ground truth: [0] | prediction: [0.0010142] |
Iter 10 | Ground truth: [1] | prediction: [0.99958826] |
Iter 11 | Ground truth: [1] | prediction: [0.99998858] |
Iter 12 | Ground truth: [1] | prediction: [0.93635657] |
Iter 13 | Ground truth: [0] | prediction: [5.21899493e-06] |
Iter 14 | Ground truth: [1] | prediction: [0.99996796] |
Iter 15 | Ground truth: [1] | prediction: [0.99998732] |
Iter 16 | Ground truth: [1] | prediction: [0.99998803] |
Iter 17 | Ground truth: [0] | prediction: [1.22098112e-05] |
Iter 18 | Ground truth: [0] | prediction: [4.76219678e-06] |
Iter 19 | Ground truth: [0] | prediction: [0.00043786] |
Iter 20 | Ground truth: [0] | prediction: [2.77989435e-05] |
Iter 21 | Ground truth: [1] | prediction: [0.99996743] |
Iter 22 | Ground truth: [0] | prediction: [4.49549571e-06] |
Iter 23 | Ground truth: [0] | prediction: [4.21171025e-06] |
Iter 24 | Ground truth: [0] | prediction: [2.60143534e-05] |
Iter 25 | Ground truth: [0] | prediction: [4.3213849e-06] |
...
Iter 98 | Ground truth: [0] | prediction: [1.16042308e-05] |
Iter 99 | Ground truth: [0] | prediction: [0.14894483] |
Iter 100 | Ground truth: [1] | prediction: [0.99949564] |
loss= 0.0014764541109010209 accuracy= [100.] %

```

XOR :

```

x2, y2 = generate_XOR_easy()
model = MLP(hidden_size = 5, activate = "sigmoid", learning_rate = 0.1)
model.train(x2, y2, epoch = 100000)
✓ 3.3s


epoch 0 loss : 0.28501257141830455
epoch 5000 loss : 0.2491668946915397
epoch 10000 loss : 0.24866809769428033
epoch 15000 loss : 0.24569798101133355
epoch 20000 loss : 0.2057230874273731
epoch 25000 loss : 0.05709604395989675
epoch 30000 loss : 0.01673334284997154
epoch 35000 loss : 0.005918474413453942
epoch 40000 loss : 0.00303345045423603
epoch 45000 loss : 0.0019160415536937942
epoch 50000 loss : 0.0013602939521317468
epoch 55000 loss : 0.0010379239924489946
epoch 60000 loss : 0.0008310455355576931
epoch 65000 loss : 0.0006885636392691007
epoch 70000 loss : 0.0005852067298738649
epoch 75000 loss : 0.0005072029378803439
epoch 80000 loss : 0.0004464710189918848
epoch 85000 loss : 0.0003979848040527405
epoch 90000 loss : 0.00035846803572880487
epoch 95000 loss : 0.0003257010223061783

```


Iter 1		Ground truth:	[0]		prediction:	[0.01659925]	
Iter 2		Ground truth:	[1]		prediction:	[0.99976639]	
Iter 3		Ground truth:	[0]		prediction:	[0.01655754]	
Iter 4		Ground truth:	[1]		prediction:	[0.99975685]	
Iter 5		Ground truth:	[0]		prediction:	[0.01652328]	
Iter 6		Ground truth:	[1]		prediction:	[0.99972007]	
Iter 7		Ground truth:	[0]		prediction:	[0.01649688]	
Iter 8		Ground truth:	[1]		prediction:	[0.99944487]	
Iter 9		Ground truth:	[0]		prediction:	[0.01647875]	
Iter 10		Ground truth:	[1]		prediction:	[0.9610959]	
Iter 11		Ground truth:	[0]		prediction:	[0.0164692]	
Iter 12		Ground truth:	[0]		prediction:	[0.01646849]	
Iter 13		Ground truth:	[1]		prediction:	[0.9651592]	
Iter 14		Ground truth:	[0]		prediction:	[0.01647683]	
Iter 15		Ground truth:	[1]		prediction:	[0.98923106]	
Iter 16		Ground truth:	[0]		prediction:	[0.01649432]	
Iter 17		Ground truth:	[1]		prediction:	[0.98859891]	
Iter 18		Ground truth:	[0]		prediction:	[0.01652105]	
Iter 19		Ground truth:	[1]		prediction:	[0.98812384]	
Iter 20		Ground truth:	[0]		prediction:	[0.01655702]	
Iter 21		Ground truth:	[1]		prediction:	[0.98789783]	
loss= 0.0002981355719449254 accuracy= [100.] %							

C. Learning curve (loss, epoch curve)

Linear :



XOR:



D. Anything you want to present

Layer 初始化時參數的選擇很重要，挑得不好會直接造成 Loss 無法下降，或是在訓練初期 Loss 降低得慢，用常態分佈選出始值的模型，會比均勻分佈挑初始值得模型訓練的效果更好。

對於沒看過的資料，像 XOR 這種複雜的分佈無法得到很高的準確度，線性模型的 testing set 就算很大，準確率也很高。

自行製作了產生更多點集的 generate_XOR_hard()，用於測試模型

```
def generate_XOR_hard(fraction = 0.1):
    inputs = []
    labels = []
    for i in range(int(1 / fraction) + 1):
        inputs.append([fraction * i, fraction * i])
        labels.append(0)
        if i == int(0.5 / fraction):
            continue
        inputs.append([fraction * i, 1 - fraction * i])
        labels.append(1)
    return np.array(inputs), np.array(labels).reshape(int(2 / fraction) + 1, 1)
```

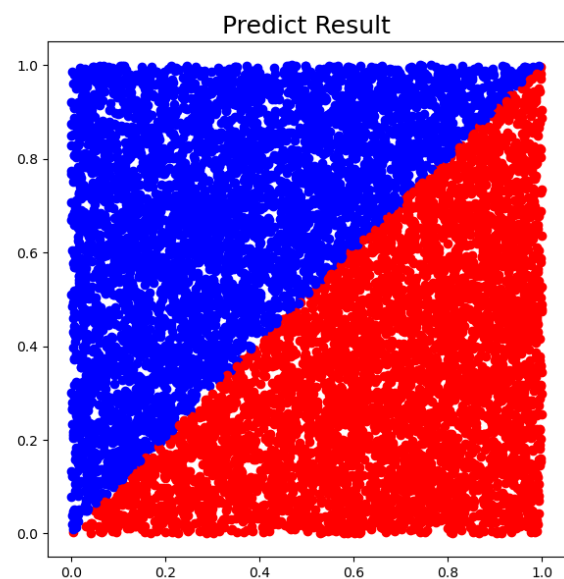

以下是測試結果：

Linear :

```



Iter 1 | Ground truth: [0] | prediction: [0.] |
Iter 2 | Ground truth: [1] | prediction: [0.99863226] |
Iter 3 | Ground truth: [0] | prediction: [0.] |
Iter 4 | Ground truth: [0] | prediction: [0.] |
Iter 5 | Ground truth: [0] | prediction: [0.] |
Iter 6 | Ground truth: [0] | prediction: [0.] |
Iter 7 | Ground truth: [1] | prediction: [1.00388766] |
Iter 8 | Ground truth: [1] | prediction: [0.99566165] |
Iter 9 | Ground truth: [0] | prediction: [0.] |
Iter 10 | Ground truth: [1] | prediction: [1.00128008] |
Iter 11 | Ground truth: [1] | prediction: [0.99831033] |
Iter 12 | Ground truth: [0] | prediction: [0.] |
Iter 13 | Ground truth: [1] | prediction: [0.99640321] |
Iter 14 | Ground truth: [0] | prediction: [0.] |
Iter 15 | Ground truth: [1] | prediction: [1.00408681] |
Iter 16 | Ground truth: [0] | prediction: [0.] |
Iter 17 | Ground truth: [0] | prediction: [0.] |
Iter 18 | Ground truth: [1] | prediction: [0.99722167] |
Iter 19 | Ground truth: [0] | prediction: [0.] |
Iter 20 | Ground truth: [0] | prediction: [0.] |
Iter 21 | Ground truth: [1] | prediction: [0.99793719] |
Iter 22 | Ground truth: [0] | prediction: [0.] |
Iter 23 | Ground truth: [0] | prediction: [0.] |
Iter 24 | Ground truth: [1] | prediction: [0.99999761] |
Iter 25 | Ground truth: [0] | prediction: [0.] |
...
Iter 9998 | Ground truth: [0] | prediction: [0.] |
Iter 9999 | Ground truth: [1] | prediction: [0.99671956] |
Iter 10000 | Ground truth: [0] | prediction: [0.] |
loss= 0.004241871682561259 accuracy= [99.36] %

```



XOR :

```
Iter 1 | Ground truth: [0] | prediction: [5.20694599e-14] |
Iter 2 | Ground truth: [1] | prediction: [1.] |
Iter 3 | Ground truth: [0] | prediction: [0.] |
Iter 4 | Ground truth: [1] | prediction: [1.] |
Iter 5 | Ground truth: [0] | prediction: [0.] |
Iter 6 | Ground truth: [1] | prediction: [1.] |
Iter 7 | Ground truth: [0] | prediction: [0.] |
Iter 8 | Ground truth: [1] | prediction: [1.] |
Iter 9 | Ground truth: [0] | prediction: [0.] |
Iter 10 | Ground truth: [1] | prediction: [1.] |
Iter 11 | Ground truth: [0] | prediction: [0.] |
Iter 12 | Ground truth: [1] | prediction: [1.] |
Iter 13 | Ground truth: [0] | prediction: [0.] |
Iter 14 | Ground truth: [1] | prediction: [1.] |
Iter 15 | Ground truth: [0] | prediction: [0.] |
Iter 16 | Ground truth: [1] | prediction: [1.] |
Iter 17 | Ground truth: [0] | prediction: [0.] |
Iter 18 | Ground truth: [1] | prediction: [1.] |
Iter 19 | Ground truth: [0] | prediction: [0.] |
Iter 20 | Ground truth: [1] | prediction: [1.] |
Iter 21 | Ground truth: [0] | prediction: [0.] |
Iter 22 | Ground truth: [1] | prediction: [1.] |
Iter 23 | Ground truth: [0] | prediction: [0.] |
Iter 24 | Ground truth: [1] | prediction: [1.] |
Iter 25 | Ground truth: [0] | prediction: [0.] |
...
Iter 199 | Ground truth: [1] | prediction: [1.] |
Iter 200 | Ground truth: [0] | prediction: [0.] |
Iter 201 | Ground truth: [1] | prediction: [1.] |
loss= 0.028381267564463724 accuracy= [96.0199005] %
```




4. Discussion (30%)

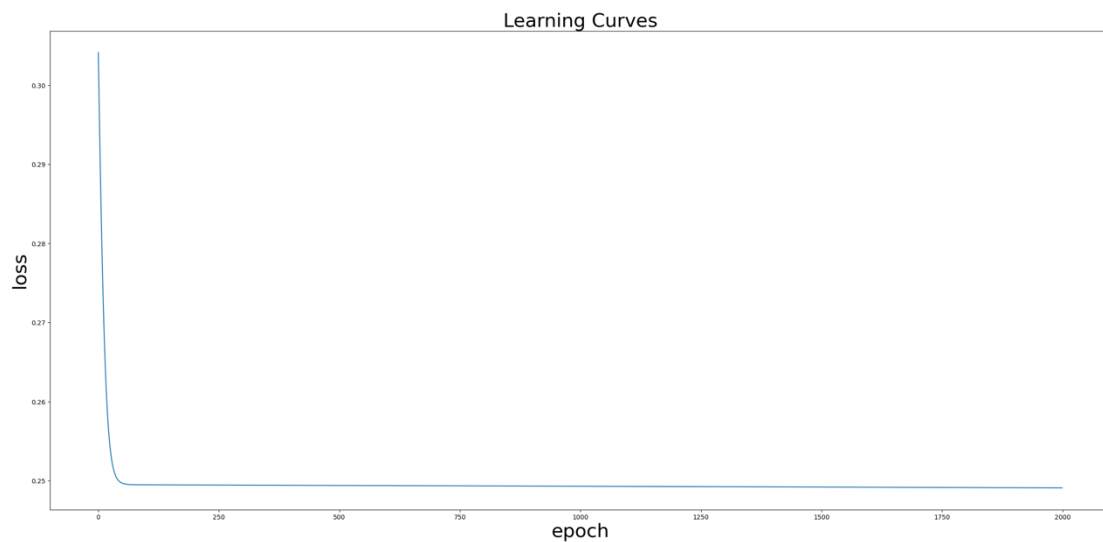
A. Try different learning rates

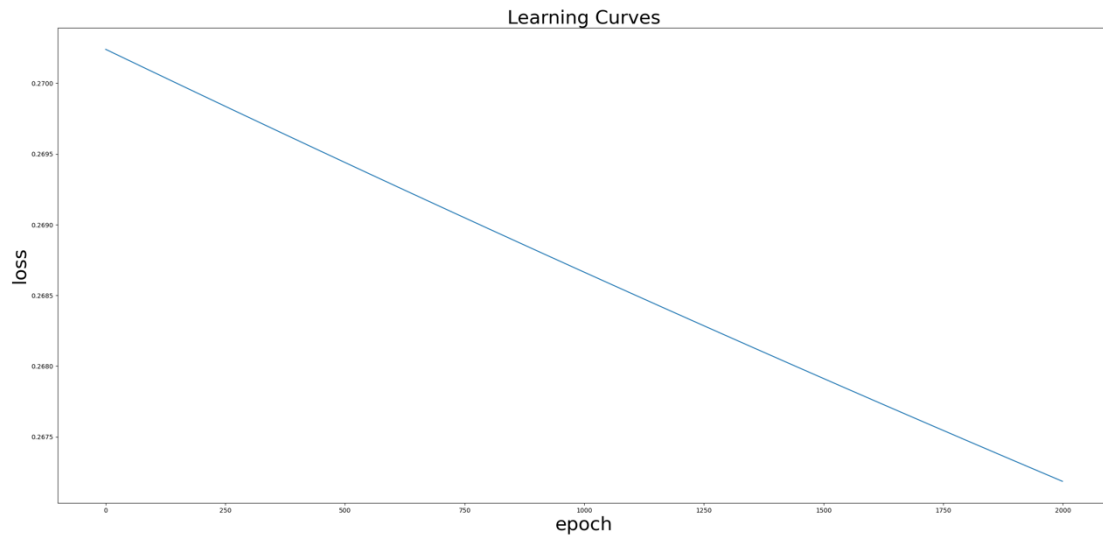
Learning rates 越小參數能進行越精確的調整，而 LR 越大能讓 Loss 收斂更快，若太大會修正過頭讓 Loss 不降反增、或是忽大忽小，但若 LR 太小的話，參數修正太少會讓訓練效果不顯著，需要更多迭代才能收斂。

LR 太大，Loss 不降反增，且前期 Loss 幾乎不變：



下面兩張圖可對比 LR 大小之差異，上方這張很陡峭(LR = 1)，下面的則是更接近一條直線(LR = 1e-3)：







B. Try different numbers of hidden units

在 linear 資料中 hidden unit size 對訓練沒有顯著影響，而在 XOR 中較大的 hidden unit 會讓訓練時間變得更長，而且 Loss function 變得複雜會產生很多 local minimum，模型比較不好收斂，若是訓練資料沒那麼複雜可以盡量減少 hidden unit size。

XOR data set, hidden unit size = 10, LR = 0.1 :



XOR data set, hidden unit size = 2, LR = 0.1 :



C. Try without activation functions


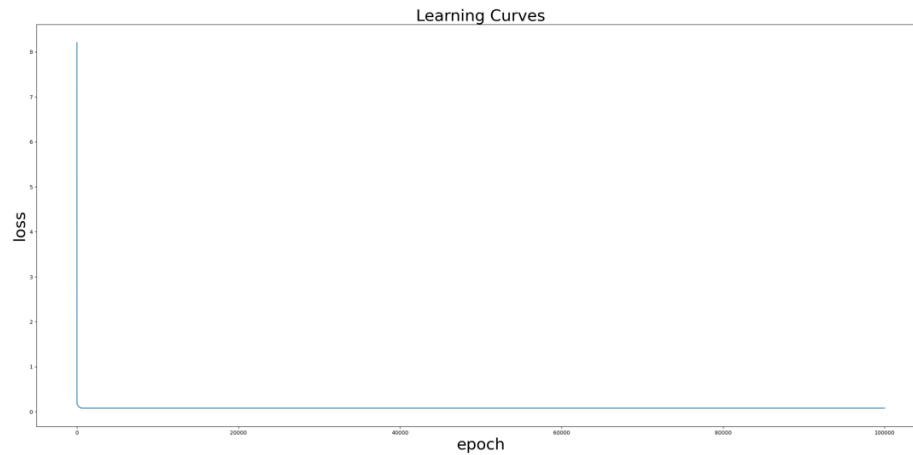
No activation functions 對 linear data set 沒顯著影響，但會導致 XOR data set 之 Loss 沒辦法收斂。

Linear data set :

```
x2, y2 = generate_linear()
model = MLP(hidden_size = 3, activate = "none", learning_rate = 1e-2)
model.train(x2, y2, epoch = 100000)
✓ 3.0s

epoch 0 loss : 8.20662186570062
epoch 5000 loss : 0.07671048204590875
epoch 10000 loss : 0.07671048204590869
epoch 15000 loss : 0.07671048204590869
epoch 20000 loss : 0.07671048204590869
epoch 25000 loss : 0.07671048204590869
epoch 30000 loss : 0.07671048204590869
epoch 35000 loss : 0.07671048204590869
epoch 40000 loss : 0.07671048204590869
epoch 45000 loss : 0.07671048204590869
epoch 50000 loss : 0.07671048204590869
epoch 55000 loss : 0.07671048204590869
epoch 60000 loss : 0.07671048204590869
epoch 65000 loss : 0.07671048204590869
epoch 70000 loss : 0.07671048204590869
epoch 75000 loss : 0.07671048204590869
epoch 80000 loss : 0.07671048204590869
epoch 85000 loss : 0.07671048204590869
epoch 90000 loss : 0.07671048204590869
epoch 95000 loss : 0.07671048204590869
```

```
Iter 1 | Ground truth: [1] | prediction: [0.8335933] |
Iter 2 | Ground truth: [0] | prediction: [0.43895091] |
Iter 3 | Ground truth: [1] | prediction: [1.29271098] |
Iter 4 | Ground truth: [0] | prediction: [-0.37490809] |
Iter 5 | Ground truth: [0] | prediction: [-0.14062501] |
Iter 6 | Ground truth: [0] | prediction: [0.44198193] |
Iter 7 | Ground truth: [0] | prediction: [0.00292047] |
Iter 8 | Ground truth: [1] | prediction: [0.59003549] |
Iter 9 | Ground truth: [0] | prediction: [0.41623762] |
Iter 10 | Ground truth: [1] | prediction: [0.57916181] |
Iter 11 | Ground truth: [1] | prediction: [1.18522516] |
Iter 12 | Ground truth: [1] | prediction: [0.81457492] |
Iter 13 | Ground truth: [0] | prediction: [0.45146687] |
Iter 14 | Ground truth: [0] | prediction: [0.01406772] |
Iter 15 | Ground truth: [1] | prediction: [0.60958108] |
Iter 16 | Ground truth: [1] | prediction: [1.07461808] |
Iter 17 | Ground truth: [1] | prediction: [0.54477349] |
Iter 18 | Ground truth: [1] | prediction: [0.80770737] |
Iter 19 | Ground truth: [1] | prediction: [1.16133821] |
Iter 20 | Ground truth: [1] | prediction: [0.67034607] |
Iter 21 | Ground truth: [1] | prediction: [1.10986141] |
Iter 22 | Ground truth: [0] | prediction: [0.16858904] |
Iter 23 | Ground truth: [1] | prediction: [0.9434397] |
Iter 24 | Ground truth: [1] | prediction: [1.18411967] |
Iter 25 | Ground truth: [0] | prediction: [0.02986747] |
...
Iter 98 | Ground truth: [0] | prediction: [0.15470842] |
Iter 99 | Ground truth: [1] | prediction: [0.6451482] |
Iter 100 | Ground truth: [0] | prediction: [-0.27126481] |
loss= 0.07671048204590869 accuracy= [100.] %
```



XOR data set :




```
x2, y2 = generate_XOR_easy()
model = MLP(hidden_size = 2, activate = "none", learning_rate = 1e-4)
model.train(x2, y2, epoch = 100000)
```

✓ 2.7s

```
epoch 0 loss : 18.499734411140444
epoch 5000 loss : 0.3005989787344996
epoch 10000 loss : 0.2785285295670828
epoch 15000 loss : 0.2685691535110954
epoch 20000 loss : 0.2630781768052016
epoch 25000 loss : 0.2596718110630188
epoch 30000 loss : 0.25738894797287404
epoch 35000 loss : 0.2557740854886236
epoch 40000 loss : 0.25458566121275894
epoch 45000 loss : 0.25368446296231034
epoch 50000 loss : 0.25298494918095016
epoch 55000 loss : 0.25243180895407225
epoch 60000 loss : 0.25198776707452997
epoch 65000 loss : 0.2516268301227043
epoch 70000 loss : 0.2513303473705192
epoch 75000 loss : 0.2510846109589962
epoch 80000 loss : 0.2508793382075476
epoch 85000 loss : 0.25070668019798115
epoch 90000 loss : 0.25056055549856976
epoch 95000 loss : 0.25043619103240233
```

Iter 1		Ground truth: [0]		prediction: [0.4998498]	
Iter 2		Ground truth: [1]		prediction: [0.54280804]	
Iter 3		Ground truth: [0]		prediction: [0.4960906]	
Iter 4		Ground truth: [1]		prediction: [0.53045719]	
Iter 5		Ground truth: [0]		prediction: [0.4923314]	
Iter 6		Ground truth: [1]		prediction: [0.51810634]	
Iter 7		Ground truth: [0]		prediction: [0.4885722]	
Iter 8		Ground truth: [1]		prediction: [0.5057555]	
Iter 9		Ground truth: [0]		prediction: [0.484813]	
Iter 10		Ground truth: [1]		prediction: [0.49340465]	
Iter 11		Ground truth: [0]		prediction: [0.4810538]	
Iter 12		Ground truth: [0]		prediction: [0.4772946]	
Iter 13		Ground truth: [1]		prediction: [0.46870295]	
Iter 14		Ground truth: [0]		prediction: [0.4735354]	
Iter 15		Ground truth: [1]		prediction: [0.4563521]	
Iter 16		Ground truth: [0]		prediction: [0.4697762]	
Iter 17		Ground truth: [1]		prediction: [0.44400125]	
Iter 18		Ground truth: [0]		prediction: [0.466017]	
Iter 19		Ground truth: [1]		prediction: [0.4316504]	
Iter 20		Ground truth: [0]		prediction: [0.4622578]	
Iter 21		Ground truth: [1]		prediction: [0.41929956]	

loss= 0.2503298182594262 accuracy= [71.42857143] %



5. Extra (10%)

B. Implement different activation functions. (3%)

實作了 ReLU，放在模型中可以得到更低的 Loss：

```
def ReLU(x):  
    return np.where(x > 0, x, 0)  
  
def derivative_ReLU(x):  
    return np.where(x > 0, 1, 0)
```

實驗結果：

```
x2, y2 = generate_linear()  
model = MLP(hidden_size = 3, activate = "ReLU", learning_rate = 0.1)  
model.train(x2, y2, epoch = 100000)
```

✓ 4.5s

```
epoch 0   loss : 0.24411503977137383  
epoch 5000 loss : 0.0015091766636624881  
epoch 10000 loss : 0.00014231161130480563  
epoch 15000 loss : 7.992985870763848e-06  
epoch 20000 loss : 3.362061168853767e-07  
epoch 25000 loss : 1.3194666365675058e-08  
epoch 30000 loss : 5.134289847974473e-10  
epoch 35000 loss : 1.97769465823814e-11  
epoch 40000 loss : 7.624948441579436e-13  
epoch 45000 loss : 2.9849661007251094e-14  
epoch 50000 loss : 1.14764727083931e-15  
epoch 55000 loss : 4.4204705359551676e-17  
epoch 60000 loss : 1.7168152847295326e-18  
epoch 65000 loss : 6.668719662989465e-20  
epoch 70000 loss : 2.563482806332538e-21  
epoch 75000 loss : 9.877692680566904e-23  
epoch 80000 loss : 3.857267875159889e-24  
epoch 85000 loss : 1.4849860415291742e-25  
epoch 90000 loss : 5.735380054599444e-27  
epoch 95000 loss : 8.26150113691776e-28
```

```
...  
Iter 98 |   Ground truth:  [1] | prediction:  [1.] |  
Iter 99 |   Ground truth:  [0] | prediction:  [0.] |  
Iter 100 |   Ground truth:  [1] | prediction:  [1.] |  
loss= 5.060446564569967e-28  accuracy= [100.] %
```

```
x2, y2 = generate_XOR_easy()
model = MLP(hidden_size = 5, activate = "ReLU", learning_rate = 0.01)
model.train(x2, y2, epoch = 100000)
```

```
epoch 0   loss : 2.777253850574109
epoch 5000 loss : 0.04435611895868108
epoch 10000 loss : 0.029009299040473478
epoch 15000 loss : 0.01922600850829182
epoch 20000 loss : 0.01079963190727733
epoch 25000 loss : 0.003328773687635289
epoch 30000 loss : 0.0005179445354492696
epoch 35000 loss : 5.4792939893598795e-05
epoch 40000 loss : 1.8500946222901952e-06
epoch 45000 loss : 2.5713404305531004e-08
epoch 50000 loss : 3.6251884289541887e-10
epoch 55000 loss : 4.995154407491181e-12
epoch 60000 loss : 6.811561241274081e-14
epoch 65000 loss : 9.647461334321907e-16
epoch 70000 loss : 1.3330858350417206e-17
epoch 75000 loss : 1.8527680946696747e-19
epoch 80000 loss : 2.589896886945452e-21
epoch 85000 loss : 3.599776512568734e-23
epoch 90000 loss : 4.986596781586126e-25
epoch 95000 loss : 8.843069117769322e-27
```

```
Iter 1 | Ground truth: [0] | prediction: [0.] |
Iter 2 | Ground truth: [1] | prediction: [1.] |
Iter 3 | Ground truth: [0] | prediction: [0.] |
Iter 4 | Ground truth: [1] | prediction: [1.] |
Iter 5 | Ground truth: [0] | prediction: [6.99440506e-15] |
Iter 6 | Ground truth: [1] | prediction: [1.] |
Iter 7 | Ground truth: [0] | prediction: [2.0539126e-14] |
Iter 8 | Ground truth: [1] | prediction: [1.] |
Iter 9 | Ground truth: [0] | prediction: [3.51940699e-14] |
Iter 10 | Ground truth: [1] | prediction: [1.] |
Iter 11 | Ground truth: [0] | prediction: [4.92939023e-14] |
Iter 12 | Ground truth: [0] | prediction: [6.29496455e-14] |
Iter 13 | Ground truth: [1] | prediction: [1.] |
Iter 14 | Ground truth: [0] | prediction: [0.] |
Iter 15 | Ground truth: [1] | prediction: [1.] |
Iter 16 | Ground truth: [0] | prediction: [0.] |
Iter 17 | Ground truth: [1] | prediction: [1.] |
Iter 18 | Ground truth: [0] | prediction: [0.] |
Iter 19 | Ground truth: [1] | prediction: [1.] |
Iter 20 | Ground truth: [0] | prediction: [0.] |
Iter 21 | Ground truth: [1] | prediction: [1.] |
loss= 1.2763376189923517e-27 accuracy= [100.] %
```